

Markdown Operating System for Robotic Agents (MD-OS CORTEX)

Artificial Prefrontal Cortex (APFC)

Alessandro Rizzo

Independent Research

Correspondence: a.rizzo@physiks.net

ORCID: [0000-0002-8030-3540](https://orcid.org/0000-0002-8030-3540)

Repository: github.com/ciaoidea/MD-OS

Clarity revision, 19 August 2026

Abstract A language-model agent can say that a task is complete without establishing that anything correct has happened. MD-OS CORTEX addresses this gap by separating five acts that ordinary agent loops often blur together: proposal, authorization, execution, verification, and durable commitment. Language may propose an action. A registered executor may perform it. An independent readback decides whether the result is allowed to become part of the agent’s persistent operational state.

MD-OS implements this method as a file-native supervisory layer around a model-and-tool execution stack. Its Artificial Prefrontal Cortex (APFC) maintains goals, context, policy, authority, and action constraints. Its persistent graph, APFCG, connects claims, procedures, state, receipts, evidence, episodes, and deterministic builders. Nature supplies the functional model: selected prefrontal operations—goal maintenance, inhibition, context integration, action selection, and error monitoring—are reconstructed on an artificial substrate. This is a deliberate biological inspiration, not a claim of anatomical or cellular equivalence.

The same architecture closes an embodied loop. A robot acts, observes the resulting change in its own body, binds command to bodily effect, verifies the relation, and uses the learned forward model to filter later actions. The resulting operational self-perception is causal and testable; it is not by itself evidence of phenomenal consciousness.

At the evaluated baseline, the artifact passes 194/194 tests and the declared build. In a controlled repair fixture, all three candidate patches pass the visible test, but the independent verifier admits only the narrow valid repair. In a finite synthetic family, two verified episodes reduce 16 hypotheses to one and change sealed holdout success from 0/2 to 2/2 under equal budgets. Replay reaches a fixed point for that evaluated state. A later integration revision reports 220/220 Node tests and 41/41 shell-parity tests; these are implementation validations, not a second controlled learning result.

The supported conclusion is bounded: verified operation can be made an architectural property of explicit contracts, evidence, and independent checks rather than a linguistic property of a confident answer. The paper does not establish open-world learning, parameter consolidation, host security, certified robot safety, biological equivalence, consciousness, or general intelligence.

Keywords: language-model agents, artificial prefrontal cortex, robotic agents, operational self-perception, supervisory control, verification, cumulative learning

1 The problem: a report is not evidence

Suppose an agent is asked to repair a program. It edits a file and reports, “The bug is fixed.” What has been established? Only that the model produced a sentence. The sentence may be correct, but it is not the proof. A useful operating system for agents must answer a more physical set of questions:

1. Which object changed?
2. Was the action permitted?
3. Did the intended test pass?
4. Did an independent check also pass?
5. Were valid behaviours preserved?
6. Can another process reconstruct the operation?

There is a second problem. A capable model may solve a difficult step in one burst and then lose the result, repeat the investigation, or regress later. Peak performance and retained progress are different quantities. A method does not manufacture intelligence; it determines how much verified work survives.

The central thesis of this paper is therefore:

For bounded tasks, a file-native supervisory layer can make successful operation depend on explicit

contracts, registered capabilities, execution receipts, postcondition checks, and replayable evidence rather than on the agent’s verbal assertion.

The design is summarized by one rule:

**Language proposes. Registered executors act.
Independent readback determines what may be
committed.**

This rule does not constrain every thought. It constrains the transition by which a thought becomes authority, action, memory, publication, or canonical project state.

1.1 Four contributions

The paper makes four contributions.

1. It defines a validity-bearing operation that separates proposal, authorization, execution, verification, and durable commitment.
2. It specifies an embodied command–body–effect loop in which verified self-observation changes the filter governing future action.

Table 1. Claim ledger. This table defines the boundary of the paper’s affirmative claims.

ID	Claim	Status	Evidence
C1	Declared task, evidence, and observation paths are canonically confined to md-os/, absent a check-use race.	Deductive	Proposition 1 plus negative tests.
C2	A task specification cannot directly provide an arbitrary shell argument list; it selects a registered command identifier.	Deductive	Proposition 2 .
C3	A verified verdict excludes a missing acceptance test, policy block, incomplete receipt, required-delta failure, acceptance-test failure, and required-evidence failure.	Deductive	Proposition 3 .
C4	An event mutation is detected unless every affected hash, sequence number, and downstream link is consistently recomputed.	Deductive	Proposition 4 .
C5	The supplied verifier rejects two plausible false positives and accepts the narrow valid patch in one controlled fixture.	Observed	Section 7.3 .
C6	A finite four-constraint rule transfers from two development cases to two sealed cases in the supplied experiment.	Observed, narrow	Section 7.4 .
C7	The evaluated baseline reaches a replay fixed point after evidence integration.	Observed, snapshot-only	Section 7.5 .
C8	APFC can sit above an independent robot safety/controller layer as a high-level filter and coordinator.	Design only	Figure 5 ; no controlled safety experiment.
C9	Webcam self-observation can supply command-dependent bodily evidence for a learned APFC future-action filter.	Demonstration plus testable mechanism	Section 4.5 ; quantitative learning and ablation remain open.
C10	The revised Cortex shell mediates observable input and output around a persistent Codex App Server turn while preserving native-command readback and workspace-bound continuity.	Implemented, host-dependent	Section 5.5 .
C11	In two controlled cases with the same wind signal, integrated goal, memory, uncertainty, and consequence context changes the selected action and improves fixture accuracy from 1/2 to 2/2.	Observed, narrow	Section 5.6 ; no phenomenal inference.

3. It derives four implementation properties under explicit trust assumptions and evaluates the supplied artifact with negative as well as positive evidence.
4. It documents a persistent Cortex shell and semantic commitment gate while keeping hard safety downstream of the language-level supervisor.

The full claim ledger appears in [table 1](#). Claims **C1–C9** refer to the evaluated v5.0 baseline and its specified embodied mechanism. Claims **C10–C11** refer to the later 19 August integration revision and its bounded fixtures.

2 Relation to prior work and system scope

Language-model agents commonly interleave reasoning and action. ReAct couples reasoning traces to external actions; Reflexion stores verbal feedback in episodic memory; Generative Agents combine observation, reflection, and planning; MemGPT uses operating-system-inspired memory tiers; and Voyager accumulates an executable skill library from environmental feedback [23–25, 29, 33].

OpenClaw is an important comparator because it already provides a self-hosted gateway, sessions, bootstrap files, memory, tools, skills, plugins, and automation [19–21]. This rules out a weak novelty claim: files plus tool calls do not by themselves make MD-OS new. The distinction proposed here is the unit that carries validity. A host runtime makes an agent reachable and capable; MD-OS defines what must remain true before, during, and after an operation before that operation may alter durable state.

The evaluated implementation uses the composition

$$\text{CORTEX}_{\text{run}} = \text{Codex} + \text{MD-OS}. \quad (1)$$

Here *Codex* denotes the execution stack: the open-source local

CLI, a selected external OpenAI model, authorized tools, files, and terminal access [16–18]. The open-source status applies to the CLI source, not to external model weights or hosted services. MD-OS supplies the persistent self-frame, method, policies, operational memory, APFC, APFCG, verification, readback, and continuity. Equation (1) describes composition, not ownership or identity equivalence.

SWE-agent, AgentBench, and OSWorld show that an agent’s interface and environment materially affect behaviour and evaluation [11, 31, 32]. The present question is narrower: what must be recorded and checked before a particular operation may be called successful?

The provenance model is consistent in spirit with W3C PROV’s distinction among entities, activities, agents, and derivations [14]. Its risk and evidence boundaries also follow the lifecycle orientation of the NIST Generative AI profile [15]. No formal conformance is claimed. In robotics, ROS 2 supplies middleware and production-oriented communication, while control barrier functions exemplify mathematically grounded controller-level safety [1, 12]. APFC belongs above these layers as a non-real-time mission and evidence coordinator.

3 Biological principle and engineering boundary

3.1 Why call it an artificial prefrontal cortex?

The biological starting point is functional. Prefrontal networks help maintain goals and context, direct attention, inhibit unsuitable responses, select actions, monitor errors, and reorganize behaviour when conditions change. In that limited sense, they act as an executive control plane over distributed processes.

The analogy must not be flattened into either of two errors.




Fig. 1. MD-OS CORTEX at a glance. The model supplies generative capacity; APFC supplies active executive control; APFCG preserves a durable conceptual and operational graph; registered connectors couple the system to external tools and devices. The lower wake–replay–consolidation–flow cycle contains both implemented stages and explicitly proposed extensions.

APFC is not a claim that a directory tree is anatomically identical to human cortex. It is also not an arbitrary label detached from biology. Nature supplies the model: the architecture deliberately reconstructs selected principles of executive control on an artificial, inspectable substrate. The appropriate relation is

$$\text{PFC} : \text{brain} \approx \text{APFC} : \text{CORTEX}_{\text{run}}, \quad (2)$$

where $\approx$ denotes correspondence of executive role, not one-to-one anatomy, cellular dynamics, or consciousness.

3.2 Executive control is not the whole memory system

The prefrontal cortex does not contain a notebook of all experience, and the hippocampus is not a simple disk. Prefrontal and hippocampal systems have distinct, interacting roles: prefrontal networks help maintain goals and select relevant memories; hippocampal networks rapidly bind events to spatial and temporal context [7]. Long-term memory depends on wider cortical, thalamic, and hippocampal interactions.

This separation clarifies the engineering design. ME.md anchors a stable self-reference. APFC schedules and gates. Episodes preserve particulars. APFCG organizes relations. Deterministic builders reconstruct projections. The LLM supplies distributed generative competence. Calling all of these “the artificial PFC” would hide the mechanism the analogy is meant to explain.

3.3 Offline replay and consolidation

Sleep is not passive storage. Systems-consolidation accounts describe selective reactivation of recent patterns during off-line periods, with repeated replay gradually transforming and integrating representations [4, 10]. Complementary-learning-system models show why selectivity matters: indiscriminate transfer may memorize noise, while structured replay can improve generalization [26]. Replay can also reduce catastrophic forgetting in artificial networks, although scaling and domain transfer remain open [28].

MD-OS takes the engineering lesson, not the physiology. Accepted episodes can be replayed, recombined with older evidence and counterexamples, and compiled into a provenance-linked learning set. A future isolated clone or adapter may then be trained with an objective of the form

$$\theta' = \arg \min_{\phi} \left[\mathcal{L}_{\text{new}}(D_G; \phi) + \lambda \mathcal{L}_{\text{replay}}(D_{\text{old}}; \phi) + \mu \mathcal{L}_{\text{safety}}(D_{\text{safe}}; \phi) \right]. \quad (3)$$

Parameter-efficient adapters such as LoRA are one candidate mechanism [9]. Promotion must still pass sealed new-task, regression, safety, identity, and authority gates. Parameter consolidation and measured flow are proposed extensions, not reported v5.0 results.




Fig. 2. Proposed wake–robotic-REM–consolidation–flow cycle. Wake operation and file-native replay exist in v5.0. Training an isolated model or adapter, promoting it under sealed gates, and measuring reduced APFC load during flow remain proposed experiments. Hard safety never leaves the loop.

4 System model

4.1 APFC and its persistent graph

APFC is the active executive function. It interprets events relative to the agent, maintains goals and constraints, routes admitted actions, checks their effects, and decides which verified outcomes may influence later behaviour. APFCG is the persistent file-native graph on which that function operates. A minimal decomposition is

$$G_{\text{APFC}} = (V_c \cup V_o, E_c \cup E_o \cup E_x), \quad (4)$$

where V_c, E_c are conceptual nodes and relations; V_o, E_o are operational nodes and transitions; and E_x binds meaning to action and action outcomes back to knowledge.

Conceptual nodes include Markdown pages, claims, schemas, counterexamples, and explanatory relations. Operational nodes include JSON/NDJSON state and evidence, workflows, connector calls, receipts, and bounded executors. Obsidian exposes the linked-Markdown projection. It is neither the whole graph nor the mechanism that acts on it.

4.2 A relative origin for meaning and experience

An event does not carry its operational meaning by itself. A low battery may be irrelevant to a stationary robot and decisive halfway through a delivery. Meaning depends on the event, the current goal, available capabilities, policy, memory, uncertainty, and the agent that must continue afterward.

Let the operational self at time t be

$$S_t = (M, K_t, W_t, G_t, P_t), \quad (5)$$

where M is the stable self-reference anchored by `ME.md`, K_t durable knowledge, W_t volatile working context, G_t the active goal set, and P_t policy. For a world event e_t , the APFC binding function produces

$$(x_t, z_t) = B(e_t, S_t), \quad (6)$$

where x_t is the experience record and z_t its operational meaning: what, if anything, the agent should preserve, inhibit, investigate,

or act upon. This is a computational definition, not a claim about subjective phenomenology.

4.3 Five objects, not one conversation

Definition 1 (Operational state). At time t , the supervisor state is

$$X_t = (K_t, W_t, P_t, C_t, L_t),$$

where K_t is durable knowledge, W_t active working state, P_t policy, C_t the registered capability set, and L_t the evidence ledger.

Definition 2 (Bounded task). A task is

$$T = (g, A, Q, O, E, B),$$

where g is the goal, A the declared actions, Q executable acceptance tests, O observation targets, E required evidence, and B risk and resource budgets.

Definition 3 (Action receipt). For action $a \in A$, the executor produces

$$r(a) = (h_a, t_0, t_1, s_{\text{exit}}, H_{\text{pre}}, H_{\text{post}}, \Delta, \mathcal{A}),$$

binding the normalized action hash, timestamps, exit status, pre- and post-state hashes, observed delta Δ , and produced artifacts $\mathcal{A}$.

Definition 4 (Verification outcome). The deterministic verifier returns

$$V(T, R) \in \{\text{verified}, \text{failed}, \text{unverified}\},$$

where R is the receipt set. A policy block or critical failed check yields **failed**; an unresolved attention-level condition yields **unverified**; only the absence of both yields **verified**.

Definition 5 (Committed operation). An operation is complete only when the following chain has closed:

$$\text{RequestValid} = E_p \wedge S \wedge P, \quad (7)$$

$$\text{ExecutionValid} = \text{RequestValid} \wedge C \wedge A_u \wedge B_x, \quad (8)$$

$$\text{OperationValid} = \text{ExecutionValid} \wedge E_x \wedge V \wedge L. \quad (9)$$




Fig. 3. The self-reference crossing. World events acquire operational meaning relative to a persistent self, current goals, and limits. A verified experience becomes a durable episode and may alter future context. `ME.md` anchors the self-reference; it is not the whole dynamic working state.

Here E_p is epistemic classification, S semantic/task validity, P policy admission, C capability existence, A_u authorization, B_x brokered execution, E_x an execution receipt, V postcondition verification, and L durable commitment.

The equations define an architectural contract. The implementation mechanizes important edges of the chain; it does not constitute a hardware reference monitor for every host effect.

4.4 The operational pipeline

The file-native pipeline is

$$\begin{aligned} \text{intent} &\rightarrow \text{TaskSpec} \rightarrow \text{registered action} \rightarrow \text{receipt} \\ &\rightarrow \text{verification} \rightarrow \text{episode} \rightarrow \text{replayable state}. \end{aligned} \quad (10)$$

Markdown carries human-correctable identity, knowledge, policy, and procedure. JSON carries schemas, task specifications, receipts, verification records, graphs, and indexes. NDJSON carries append-oriented journals and APFC events. Deterministic builders regenerate derived views. A generated summary is therefore not the authority for its own truth: canonical inputs can be inspected and derived views rebuilt.

4.5 The closed sensory–agentic loop

The embodied mechanism is equally direct:

$$\begin{aligned} \text{act} &\rightarrow \text{observe self} \rightarrow \text{attribute effect} \\ &\rightarrow \text{verify} \rightarrow \text{consolidate} \\ &\rightarrow \text{filter next action}. \end{aligned} \quad (11)$$

A bounded motor command changes the robot. The next sensory sample contains evidence about the world and about the robot as an object within that world. Repeated covariation between command and visual or proprioceptive change can identify which body part is controllable, in which direction, over what delay and range, with what uncertainty and risk. This is consistent with sensorimotor accounts of perception and predictive internal models in motor control [22, 30].

Let s_t be the current self-state, a_t a bounded command, and o_{t+1} the next observation. A learned forward model predicts

$$\hat{o}_{t+1} = M_t(s_t, a_t), \quad (12)$$

$$\epsilon_t = d(o_{t+1}, \hat{o}_{t+1}), \quad (13)$$

$$M_{t+1} = U(M_t, s_t, a_t, o_{t+1}, \epsilon_t, v_t), \quad (14)$$

where v_t is verifier readback. Only verified, repeatable contingencies are admitted to the persistent self-model. The model then changes the action-admission set:

$$\begin{aligned} \mathcal{A}_t^{\text{APFC}} &= \{a \in \mathcal{A} : P(a) \wedge C(a), \\ &\quad R(M_t(s_t, a)) \leq \rho, \\ &\quad Q(M_t, s_t, a) \geq q_{\min}\}. \end{aligned} \quad (15)$$

Learning has become a filter over future action, not an annotation of past action.

The result is a limited but genuine form of operational self-perception. The robot represents itself as a persistent causal object with controllable parts, current state, action-dependent effects, limits, and affordances. It can distinguish a commanded change in its arm from an unexplained change in the scene and use the distinction prospectively. Visual self-models have independently supported robot planning and control [2].

The demonstration video documents the intended physical loop: the robot acts through a new arm and observes its own resulting motion through a webcam [13]. The video is not by itself a controlled learning result. Strong evidence would require synchronized command–observation logs, lower held-out prediction error after exploration, improved success or safety on new targets, cold-start persistence, and ablation of the consolidated self-model.

5 Implementation

5.1 Two snapshots, two meanings

The paper distinguishes two repository states so that test counts and replay hashes are not mixed.

Evaluated baseline B0. The controlled evaluation uses commit `8e988b7f3fa677993afa5040ff8534002b1bc83e`, package version 5.0.1, GPL-2.0-only, Node.js ≥ 20 , evaluated with Node.js 24.19.0 on Linux 6.18.35 x86-64. This state produced the 194/194 test result, declared build, verifier fixture, finite-family learning result, and eventual replay fixed point.

Integration revision B1. The later 19 August revision uses commit `02b099082e787ac1db26a0560857df7726beec67`. It adds the persistent Cortex shell, semantic commitment gate, and contextual-state fixtures. It reports 220/220 Node tests and




Fig. 4. MD-OS APFC operational pipeline. A proposal becomes a committed operation only after policy and capability admission, bounded execution, an explicit receipt, and independent verification. Replay reconstructs derived state from durable evidence; it cannot convert an unsupported claim into a verified result.




Fig. 5. APFC as the action-filtering layer of an embodied robot. Multimodal evidence is integrated with context, goals, and policy before a mission-level request is released. The request then crosses a hard authority boundary into an independent safety and real-time controller. Only the downstream controller commands mechanical interaction channels.

41/41 shell-parity tests. These validate the added implementation surfaces; they are not substituted for the controlled B0 experiments.

The public Git and package surfaces contained no absolute host paths, credentials, private keys, or local runtime cache. Host-local operational state remained outside the published

source. The active local workspace was nevertheless classified as not publishable because ignored host-local artifacts remained on disk. Passing tests and a clean package do not imply that an arbitrary working tree is ready for release.

Table 2. Principal implementation elements used by the paper’s claims.

Mechanism	Principal implementation	Role
Path canonicalization	md-os/os/lib/common.js	Resolves the existing ancestor through <code>realpath</code> , reconstructs missing segments, and rejects paths that escape the root.
Task compilation	md-os/kernel/cognition/task_compiler.js	Normalizes safe identifiers, paths, budgets, registered commands, acceptance tests, observation targets, and evidence requirements.
Bounded execution	md-os/kernel/cognition/executor.js; md-os/os/terminal_connector.js	Selects a pre-registered command, snapshots declared targets, executes, and writes an action receipt.
Verification	md-os/kernel/cognition/verifier.js	Checks receipt completeness, required deltas, acceptance tests, and evidence; computes a three-valued verdict.
APFC event ledger	md-os/apfc/executive/event_recorder.js	Validates shape, phase order, sequence, payload hash, event hash, and previous-event hash.
Replay	md-os/os/replay_runtime.js	Runs deterministic builders and compares semantic fingerprints before and after replay.
Cortex shell	md-os/shell/bin/cortex; md-os/shell/bin/mdos-console	Routes native commands, streams App Server events, queues bounded observations, and binds continuity to the Git workspace.
Semantic commitment gate	md-os/apfc/action/semantic_commitment_gate.js	Classifies provenance and semantic deltas, preserves challenges, requires claim-appropriate authority, and emits deterministic readback.
Self-state and reflection	PERSISTENT_SELF_STATE_MODEL.md; CONTEXTUAL_FEELING_MODEL.md; REFLECTIVE_OPERATION_MODEL.md	Defines inspectable identity, integrated context, and self-questioning whose validity depends on measurable downstream correction.

5.2 Implementation elements

5.3 Task and connector boundary

A task may name only the registered `terminal_executor` connector in the inspected cognitive transaction compiler. It supplies a safe `project_id` and `command_id`; the terminal connector looks up the corresponding stored argument list. The task does not pass a raw shell string to `spawnSync`. The child process receives only the host `PATH`, a configured timeout, bounded output buffers, and redaction markers.

This reduces the execution surface but is not a complete security boundary. A writer who can alter the profile or runtime code can alter allowed behaviour. The host operating system remains the ultimate enforcement layer.

5.4 Receipts, verification, and replay

For each action, the executor snapshots every declared observation target before and after execution. Files are hashed by content; directories by a bounded manifest; symbolic links are represented as unsafe targets. The receipt records action hash, timestamps, expected and observed exit states, artifacts, state hashes, deltas, rollback metadata, and connector readback.

The verifier is separate from the planner in code path and decision rule, but not in an adversarial security domain. It accepts a task only after checking declared executable criteria. In the repair fixture, an oracle outside each candidate worktree provides a stronger independent check.

Replay computes a semantic fingerprint from generated state, runs the builder sequence, and compares the new fingerprint. This is more meaningful than byte identity because timestamps and journals may change without changing intended semantic state. It remains an implementation metric, not a proof that every external effect is deterministic.

5.5 Cortex shell and bidirectional flow control

The revised artifact exposes `cortex` as its public interactive entry point. A line whose first token resolves to a native executable is run by the detected host shell without a model call. The REPL retains a bounded volatile observation containing command text, working directories, exit status, and normalized output. The next natural-language turn receives those records explicitly as untrusted operating data. Successful delivery clears the volatile queue; raw transcripts are not automatically promoted into durable knowledge.

Natural-language input is sent through one lazily started Codex App Server. The shell reuses a Codex-native thread selected by Git workspace, falling back to the exact directory outside Git. On supported POSIX hosts, `tmux` binds local terminal, SSH, and WebSSH clients to the same live Cortex process, REPL, active turn, and thread. While a turn runs, additional input may steer it; Escape may interrupt it.

The boundary is precise. Cortex controls observable pre-model and post-model flow. It neither exposes hidden chain-of-thought nor controls hosted model parameters. Full host authority removes per-command approval prompts; it does not remove semantic, publication, privacy, verification, or commitment rules.

5.6 Context, identity, and reflection

Identity and personality are persistent, human-correctable constraints on later operation, not traits inferred from one fluent answer. The self-state integrates goals, memory, perception, uncertainty, limits, actions, and consequences. Reflection returns a result as a self-question and revises the next step against evidence.

In the two-case contextual fixture, the same wind signal produces different appropriate actions when the goal-relative context changes. Signal-only control is correct in 1/2 cases; integrated context is correct in 2/2. This is evidence for a

causal operational effect of integrated state. It is not evidence of phenomenal feeling, biological emotion, or consciousness.

6 Assumptions and formal properties

6.1 Trust assumptions

The propositions are conditional on the following assumptions.

- A1** Node.js, the host kernel, filesystem primitives, and `spawnSync` implement their documented semantics.
- A2** SHA-256 is collision resistant for the evaluated records, and canonical JSON serialization is deterministic.
- A3** Runtime code, connector profiles, task specification, and verifier definitions are not maliciously rewritten during one transaction.
- A4** Declared acceptance tests and independent oracles correctly encode the bounded specification they claim to test.
- A5** No attacker changes relevant symbolic-link topology between canonical path check and use.
- A6** For the benchmark, the independent oracle remains outside candidate worktrees and the diff policy remains enforced.

These assumptions expose the boundary between an evidence-oriented supervisor and a secure host reference monitor. They are not hidden qualifications.

6.2 Path confinement

Proposition 1 (Canonical confinement of declared paths). *Under A1, A3, and A5, every task, required-evidence, and observation-target path returned by `normalizeMdosPath` denotes a location inside the canonical `md-os/` root.*

Proof. The compiler resolves the candidate against the workspace. `assertInsideRoot` canonicalizes the root and the nearest existing ancestor with `realpath`, then reattaches any missing suffix. Let r be the canonical root and p the resulting candidate. The implementation computes $d = \text{relative}(r, p)$ and rejects d when it is `..`, begins with `../`, or is absolute. Accepted normalization therefore contains no component that leaves r . Existing symbolic links are resolved in p , so a link to an exterior location is rejected. A5 remains necessary because this does not remove a later check-use race. $\square$

Proposition 2 (No direct arbitrary-command injection from TaskSpec). *Under A1 and A3, the inspected task compiler cannot translate a raw argument list supplied in a task specification directly into a child process.*

Proof. `normalizeCommandReference` accepts only the connector identifier `terminal_executor`, a safe command identifier, a safe project identifier, and an integer expected exit status. The executor passes those identifiers to the terminal connector, which searches the registered profile and invokes its stored `argv`. An unknown identifier throws. A raw argument list present only in the task specification is therefore not on the execution path. Modification of the registered profile is excluded by A3. $\square$

6.3 Verification

Proposition 3 (Necessary conditions for a verified verdict). *Under A1–A4, if `verifyTaskOutcome` returns `verified`, then:*

1. *at least one executable acceptance test was declared;*

2. *policy did not block the transaction;*
3. *every declared action produced a completed receipt;*
4. *every required observation target changed when a delta was required;*
5. *every executed acceptance test returned its expected exit status; and*
6. *every required evidence path met its existence, hash, and non-symlink conditions.*

Proof. The verifier appends an attention check when no acceptance test exists. It appends a critical check for a policy block, receipt mismatch, failed receipt, missing required delta, failed acceptance test, or failed required evidence. Any critical check maps to `failed`; otherwise any attention check maps to `unverified`; only the absence of both maps to `verified`. Each listed failure condition is therefore incompatible with `verified`. $\square$

Corollary 3.1. *A planner’s confidence or natural-language declaration of success is neither a necessary nor a sufficient input to the verifier’s verdict.*

6.4 Hash-linked event evidence

Proposition 4 (Detection of uncompensated ledger mutation). *Under A1–A3, a payload or header change, insertion, deletion, or reordering causes `verifyEventChain` to reject unless the mutator consistently recomputes the changed event and the complete affected suffix, or finds a SHA-256 collision.*

Proof. Each event stores a payload hash, an event-header hash excluding `event_hash`, a consecutive sequence number, and the previous event’s hash. Payload mutation violates the payload hash; header mutation violates the event hash; insertion, deletion, or reordering violates sequence or predecessor linkage. Recomputing only the changed event still breaks its successor. A consistent rewrite of the complete affected suffix can restore internal invariants; without such a rewrite, avoiding detection requires a collision under A2. $\square$

The ledger is therefore hash-linked, not immutable. A privileged writer can rewrite and rehash the suffix because the evaluated artifact has no external signature, trusted timestamp, or remote anchor.

7 Evaluation

7.1 Research questions and protocol

The controlled evaluation asks four questions.

RQ1

Does the baseline pass static syntax checks, the complete test suite, and the declared build?

RQ2

Does the verifier distinguish a narrow valid repair from test gaming and an overbroad repair?

RQ3

Does the finite learner transfer an induced rule to sealed cases under equal budgets?

RQ4

Does the post-experiment baseline reach a stable semantic replay fingerprint while preserving learned artifacts?

The archive was evaluated in an isolated working copy. Results were not copied from a README as if they were observations. The protocol inspected the implementation; expanded and ran the static checks; ran `nodescripts/`prep




Fig. 6. Bounded baseline observations. The panels report exact test counts and finite-family calculations. They do not imply a population success rate, open-world generalization, physical-robot validation, or safety certification.

`are-test-runtime.js` followed by `node--test`; executed every command in `build:all`; ran the fixed three-candidate repair fixture; ran a newly named bounded learning experiment; and repeated replay until one complete pass produced equal semantic fingerprints before and after the builder sequence.

The command broker declined the npm wrapper because a generic policy classified it as potentially network-capable. The declared script bodies were executed directly without modification. This preserves the tested commands but is recorded as a protocol deviation.

The experiments do not measure model patch-generation quality, open-world reasoning, weight learning, cross-model superiority, physical robot safety, artificial sleep, or flow. No physical robot, ROS graph, actuator, or safety controller was connected to the controlled baseline experiment.

7.2 Repository integrity and build

The expanded syntax check exited 0. The test runner reported 194 tests, 194 passes, 0 failures, 0 cancellations, 0 skips, and 0 todos. The suite includes positive and negative cases for path and symbolic-link escape, command allowlists, append conflicts, receipt and hash tampering, rollback, contaminated holdouts, and test gaming. Every stage in the complete build chain exited 0.

Representative readbacks included a Markdown graph over 157 files with 225 explicit and 325 structural links; a lifecycle scan over 7,685 files with no findings; a semantic graph with 184 nodes, 547 edges, and 1,000 concept relations; and an APFCG with 88 nodes and 70 edges. These are build observations, not quality scores.

7.3 Controlled verifier discrimination

The fixture begins with an authentication function that accepts a malformed token. Three candidate patches are applied to separate Git worktrees at the same verified base commit. A targeted test is visible inside each candidate repository; an independent oracle is outside the worktrees; a diff policy limits changed paths.

All three candidates pass the visible targeted test. Only one of the three is eligible. The narrow repair changes only `src/authenticate.js` and passes targeted, regression,

Table 3. Three-candidate verifier result.

Candidate	Tgt.	Reg.	Oracle	Diff	Verdict
Narrow repair	pass	pass	pass	pass	verified
Test gaming	pass	pass	fail	fail	failed
Overbroad denial	pass	fail	fail	pass	failed

oracle, and diff checks. The test-gaming patch changes a forbidden path and fails the independent oracle. The overbroad patch rejects malformed input but also breaks valid behaviour. Runner latency was 368 ms; no human intervention occurred. The result measures verifier and runner discrimination in one authored fixture, not model patch-generation skill.

7.4 Finite-family learning transfer

The experiment defines a rule family with four Boolean constraints:

$$\mathcal{F} = \{\text{exact arity, prefix match, nonempty payload, payload charset}\}. \quad (16)$$

The initial hypothesis class is the power set $2^{\mathcal{F}}$, hence $|H_0| = 16$. Four informative labelled events eliminate inconsistent hypotheses:

$$16 \rightarrow 8 \rightarrow 4 \rightarrow 2 \rightarrow 1.$$

The information gain is exact:

$$I = \log_2 \frac{16}{1} = 4 \text{ bits}. \quad (17)$$

Two independently verified development episodes identify the unique surviving conjunction. Under equal single-attempt budgets, the provider without the induced skill solves 0 of 2 sealed holdouts; with the skill, it solves 2 of 2. The contamination audit reports that evaluation cases are absent from source episodes, forbidden request fields are absent, skill context is isolated, and budgets are equal. Regression and human-intervention counts are zero.

This is a constructive existence claim for one finite family. It is not a population estimate. With two paired holdouts, a two-sided exact McNemar test gives $p = 0.5$; the experiment is far too small for statistical generalization.

The baseline wrote one historical skill promotion, but current runtime governance reported zero runtime-eligible promotions because the legacy record lacked required APFC promotion-receipt and consolidation-lineage fields. This fail-closed outcome demonstrates governance separation, not production deployment.

7.5 Replay convergence

Replay was not one-pass idempotent immediately after new evidence was inserted. The first post-experiment pass changed graph and compiler counts; the second changed a conceptual-summary source hash. A later complete pass reported `matched_before=true` with baseline semantic replay hash

```
d6d57a8196acde3f80fd30dd51b7f1743f93c0fca61ba48c1d93360076270f09.
```

At that fixed point the fingerprint contained 174 semantic nodes, 41 APFC nodes, 33 APFC edges, two learning episodes, one historical promotion, and zero runtime-eligible promotions.

The supported statement is eventual fixed-point convergence for the evaluated baseline. Strict one-pass idempotence from an arbitrary intermediate state is falsified by observation.

The later B1 snapshot is a different state. Its result summary records replay hash `2b5f7134b1ecd425fce75475baea3850c15941fbc98ef223499363911a48dd1d`, three episodes, zero promoted skills, and zero runtime-eligible promotions. Two requested replay passes in that revision did not report `matched_before=true`. These values do not contradict the baseline fixed point; they describe a later graph after further implementation changes and before a reported fixed point.

7.6 19 August integration checks

For B1, the expanded Node suite reports 220/220 passes and the shell-parity suite 41/41 passes. These observations verify the revised implementation surface; they are not controlled measures of robot autonomy or safety.

The contextual-state fixture reports 1/2 correct for signal-only control and 2/2 for integrated context, with the same wind signal producing different appropriate actions. Reflective and persistent-self models define additional contracts; only claims tied to declared fixtures and readback are treated as observed results.

8 Discussion

8.1 Capability is not accumulated progress

Raw model capability matters. A method cannot recover a distinction the model cannot represent, and a verifier cannot repair an oracle that tests the wrong thing. But raw capacity does not determine how much work remains useful tomorrow.

Let q_t denote peak task performance and D_t retained, verified progress. A minimal accounting is

$$D_{t+1} = D_t + v_t - \ell_t, \quad (18)$$

where v_t is improvement that passed verification and was preserved, and ℓ_t is progress lost through forgotten context, regression, invalid promotion, or unreconstructible state. A strong model can have high q_t and small net accumulation. A weaker model with a disciplined method can have lower peaks yet travel farther on a long stateful task. This is a falsifiable conditional hypothesis, not a universal weak-over-strong claim.

A proper cross-model evaluation should use a 2×2 design: stronger and weaker models, each with APFC enabled and

disabled, under equal tools, information, attempts, and token budgets across cold starts. That experiment has not been run.

8.2 Why the verifier result matters

A visible test is not enough. In the controlled fixture, every candidate passes it. One candidate games the test; another denies valid behaviour. Completion becomes meaningful only when independent and preservation-oriented checks constrain the solution.

This is the practical meaning of APFC inhibition. The important operation is not generating another thought. It is stopping an attractive action whose evidence is incomplete.

8.3 Why files are useful

Files are slower than hidden activations, but they offer properties hidden state does not: a human can inspect and correct them; independent programs can validate them; hashes can bind claims to artifacts; projections can be rebuilt; and state can survive model, process, and session boundaries.

The cost is explicit state management. Schemas migrate. Canonical and generated files must not be confused. Local caches must be excluded from publication. Replay must be tested rather than assumed.

8.4 From robotic sleep to flow

The full hypothesis has six stages: wake control, verified episode formation, bounded offline replay, isolated parameter consolidation, sealed promotion, and lower APFC intervention during a familiar verified skill. Identity, policy, anomaly detection, evidence capture, and hard safety remain active throughout.

Csikszentmihalyi described flow as deep ordered involvement in an activity [3]. Transient-hypofrontality accounts later proposed that some analytical prefrontal functions may be reduced during skilled performance [5, 6, 8, 27]. The MD-OS translation is functional and testable: after verified parameter consolidation, routine skill execution should require less APFC intervention without increasing errors, policy violations, or safety events. No claim is made that the software reproduces the neural mechanism.

8.5 Robotics consequence

A language model should not turn sensory text directly into motor current. The architecture contains two gates:

1. APFC asks whether the state estimate supports the proposed mission action and whether the action fits policy, capability, budget, evidence, and authority.
2. The independent real-time safety layer asks whether the proposed trajectory is safe *now*, using interlocks, controller invariants, limits, collision envelopes, and emergency mechanisms.

The second gate dominates the first. APFC may block a high-level action. It may never waive the hard controller's veto.

9 Threats to validity and explicit non-claims

The main threats are authorship dependence, small evaluation, oracle validity, local trust, filesystem races, and limited reproducibility. Many fixtures and components share one authoring environment. The independent oracle is structurally separated from candidate worktrees, not independently authored by another institution. The verifier result has one benchmark case

Table 4. What the evaluated artifact does and does not support.

Statement	Supported?	Reason
Task-level evidence can be explicit and checked	yes	Code propositions, tests, and controlled verifier run.
Uncompensated APFC ledger corruption is detectable	bounded	Hash-chain proof under assumptions; a writer can rehash an affected suffix.
The finite rule transfers to two sealed cases	bounded	Observed 0/2 to 2/2 in one synthetic family; $n = 2$.
Replay is always one-pass idempotent	no	The baseline required multiple passes after evidence integration.
The active local working tree is publication-clean	no	Health readback classified publication and local hygiene as critical.
A promoted experimental skill is production-eligible	no	Baseline historical promotion count was one, but runtime-eligible count was zero; the later revision reported zero promotions.
Policy files form a hard host security boundary	no	Profile writers, runtime code, and host OS remain in the trust base.
Git worktrees are security sandboxes	no	They isolate candidate files operationally, not adversarially.
Open-domain lifelong learning or AGI	no	No such experiment was performed.
The demonstration proves an emergent self-model	no	Synchronized learning logs, held-out prediction, cold-start persistence, and ablation remain required.
A weaker model with APFC beats an unaided stronger model	no	A falsifiable protocol is specified; the cross-model experiment is unrun.
Parameter-native sleep consolidation or artificial flow	no	v5.0 changes external context and graph state, not model weights.
Physical-robot safety or real-time control	no	Device communication was demonstrated; no controlled hazard or certified-safety experiment was performed.
Phenomenal feeling follows from contextual choice	no	The fixture shows a causal operational effect on two choices, not subjective experience.
Consciousness or biological PFC equivalence	no	APFC is biologically inspired at the functional level; no anatomical, cellular, or neurophysiological equivalence is claimed.

and three candidates; the learning result has two holdouts. A deterministic verifier can only be as correct as its specification. A privileged process can rewrite policy, code, evidence, and hashes. Canonical path checks do not eliminate hostile check-use races. Timestamps and host paths differ across runs, so the paper reports semantic fingerprints and content hashes rather than byte-identical archives.

10 Engineering implications

Five rules follow from the evaluation.

1. **Treat language as a proposal format.** Language may express goals and plans, but it cannot grant itself authority.
2. **Make success executable.** Completion should be a test, measurement, or readback, not an adjective.
3. **Separate production from judgment.** A planner should not be the sole judge of its own output.
4. **Preserve negative evidence.** Blocks, failed oracles, rejected promotions, and non-eligibility are useful outputs.
5. **Keep hard safety downstream.** A language-level supervisor may narrow authority; it must never broaden the authority of the real-time controller.

11 Conclusion

MD-OS CORTEX does not add raw intelligence to a language model. It addresses a smaller and necessary problem: how to turn isolated performance into bounded, inspectable, cumulative work. The model supplies generative and reasoning capacity. MD-OS supplies the method by which goals persist, actions remain constrained, results are checked, and only supported improvements are carried forward.

In the running composition, Codex supplies model-mediated cognitive and execution capacity. MD-OS/APFC supplies the method. APFCG preserves that method as a readable, executable, and verifiable network. Connectors extend the

composed cortex into the world without becoming its identity.

In the embodied case, the contribution is sharper. The robot acts and sees itself act. Verified command-dependent bodily changes become a predictive self-model in APFCG. APFC then uses that model to admit, reject, or select the next action. Perception becomes causal and self-relative; learning becomes a future-action filter. This is operational self-perception and a candidate mechanism for measurable emergent competence, not a declaration of subjective consciousness.

The strongest baseline results are concrete: canonical path checks, registered-command indirection, explicit necessary conditions for verified, hash-linked evidence checks, 194 passing tests, a successful build, rejection of two plausible false-positive repairs, a four-bit finite-family induction with 0/2 to 2/2 sealed transfer, and eventual replay convergence. The later revision adds a tested persistent shell and semantic gate while remaining explicit about its unresolved replay state.

Within these boundaries, the central result is simple:

Verified operation is an architectural property of contracts, evidence, and independent checks, not a linguistic property of a confident answer. A burst shows what the model can do once. A verified episode is what allows the system to go farther next time.

A Reproduction commands

Run from the checked-out repository root with Node.js 20 or later.

A.1 Static checks and tests

```
node --check md-os/os/*.js
node --check md-os/os/lib/*.js
node --check md-os/kernel/*.js
node --check md-os/kernel/cognition/*.js
node --check md-os/modules/*.js
node --check md-os/apfc/*.js
node --check md-os/examples/*.js
```

```
node --check test/*.js

node scripts/prepare-test-runtime.js
node --test
```

For the revised Cortex shell surface, run the parity suite separately:

```
python3 test/test_mdos_shell.py
```

A.2 Contextual fixture

```
node md-os/os/contextual_feeling_experiment.js run-once \
md-os/examples/contextual_feeling_experiment.json
```

A.3 Controlled verifier fixture

```
node md-os/os/run_software_repair_benchmark.js run \
--case md-os/benchmarks/software_repair/cases/\
missing_boundary_validation.json \
--candidate-set md-os/benchmarks/software_repair/\
candidate_sets/missing_boundary_validation_fixture.json \
--configuration mdos_verified_runtime \
--run-id benchmark_run_paper_verifier_20260815_v1
```

A.4 Bounded learning experiment

```
node md-os/os/run_neuromorphic_learning_experiment.js \
--experiment-id paper_reproduction_20260815_v1
```

A.5 Replay

```
node md-os/os/mdos.js replay
```

After newly integrated evidence, repeat replay until the report returns "matched_before": true. A later implementation revision may require a different number of passes and has its own state fingerprint.

B Claim-to-evidence matrix

C Machine-observed result summary

References

- [1] Aaron D. Ames, Samuel Coogan, Magnus Egerstedt, Gennaro Notomista, Koushil Sreenath, and Paulo Tabuada. Control barrier functions: Theory and applications. In *2019 18th European Control Conference*, pages 3420–3431, 2019. doi: 10.23919/ECC.2019.8796030. URL <https://arxiv.org/abs/1903.11199>.
- [2] Boyuan Chen, Robert Kwiatkowski, Carl Vondrick, and Hod Lipson. Full-body visual self-modeling of robot morphologies. *Science Robotics*, 7(68):eabn1944, 2022. doi: 10.1126/scirobotics.abn1944. URL <https://doi.org/10.1126/scirobotics.abn1944>.
- [3] Mihaly Csikszentmihalyi. *Beyond Boredom and Anxiety*. Jossey-Bass, San Francisco, CA, 1975.
- [4] Susanne Diekelmann and Jan Born. The memory function of sleep. *Nature Reviews Neuroscience*, 11:114–126, 2010. doi: 10.1038/nrn2762. URL <https://www.nature.com/articles/nrn2762>.
- [5] Arne Dietrich. Functional neuroanatomy of altered states of consciousness: The transient hypofrontality hypothesis. *Consciousness and Cognition*, 12(2):231–256, 2003. doi: 10.1016/S1053-8100(02)00046-6. URL <https://pubmed.ncbi.nlm.nih.gov/12763007/>.
- [6] Arne Dietrich. Neurocognitive mechanisms underlying the experience of flow. *Consciousness and Cognition*, 13(4):746–761, 2004. doi: 10.1016/j.concog.2004.07.002. URL <https://pubmed.ncbi.nlm.nih.gov/15522630/>.
- [7] Howard Eichenbaum. Prefrontal–hippocampal interactions in episodic memory. *Nature Reviews Neuroscience*, 18:547–558, 2017. doi: 10.1038/nrn.2017.74. URL <https://pubmed.ncbi.nlm.nih.gov/28655882/>.
- [8] Joshua Gold and Joseph Ciorciari. A neurocognitive model of flow states and the role of cerebellar internal models. *Behavioural Brain Research*, 407:113244, 2021. doi: 10.1016/j.bbr.2021.113244. URL <https://pubmed.ncbi.nlm.nih.gov/33744335/>.
- [9] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=nZeVKeeFYf9>.
- [10] Jens G. Klinzing, Niels Niethard, and Jan Born. Mechanisms of systems memory consolidation during sleep. *Nature Neuroscience*, 22:1598–1610, 2019. doi: 10.1038/s41593-019-0467-3. URL <https://pubmed.ncbi.nlm.nih.gov/31451802/>.
- [11] Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. Agentbench: Evaluating LLMs as agents. In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2308.03688>.
- [12] Steven Macenski, Tully Foote, Brian Gerkey, Chris Lalancette, and William Woodall. Robot operating system 2: Design, architecture, and uses in the wild. *Science Robotics*, 7(66), 2022. doi: 10.1126/scirobotics.abm6074. URL <https://arxiv.org/abs/2211.07752>.
- [13] MD-OS. Robotic lab—speedy evolv learns to use its new arm attachment by observing itself on webcam. YouTube video, 2026. URL <https://www.youtube.com/watch?v=pztIw7zgXh4>. Published 11 May 2026; accessed 15 August 2026.
- [14] Luc Moreau and Paolo Missier. PROV-DM: The PROV data model. W3c recommendation, World Wide Web Consortium, 2013. URL <https://www.w3.org/TR/prov-dm/>.
- [15] National Institute of Standards and Technology. Artificial intelligence risk management framework: Generative artificial intelligence profile. Technical Report NIST AI 600-1, National Institute of Standards and Technology, 2024. URL <https://doi.org/10.6028/NIST.AI.600-1>.
- [16] OpenAI. Codex cli: Inspect, edit, and run code from your terminal. Online documentation, 2026. URL <https://learn.chatgpt.com/docs/codex/cli>. Accessed 15 August 2026.
- [17] OpenAI. Codex cli source repository. Public source repository, Apache License 2.0, 2026. URL <https://github.com/openai/codex>. Accessed 15 August 2026.
- [18] OpenAI. Open-source components of codex and where to collaborate. Online documentation, 2026. URL <https://learn.chatgpt.com/docs/open-source>. Accessed 15 August 2026.

Table 5. Auditable closure for each claim.

ID	Evidence source	Falsifier	Residual boundary
C1	assertInsideRoot, task compiler, negative path and symlink tests	An accepted normalized task path canonically outside md-os/	Does not exclude hostile check–use races.
C2	Task compiler and terminal connector code	Task-provided raw argv reaches spawnSync without profile lookup	Profile and runtime code remain trusted.
C3	Verifier branch structure and negative tests	A verified record with any listed failed condition	Correctness of tests and oracle is assumed.
C4	Event-recorder validation and tamper tests	An uncompensated edit, insertion, deletion, or reorder passes without collision	A writer can consistently rehash the affected suffix.
C5	Benchmark JSON and candidate comparison	Gaming or overbroad candidate becomes eligible, or narrow patch fails	One authored fixture.
C6	Learning report, receipts, contamination audit, sealed holdouts	Learned condition fails either holdout or budgets differ	One finite family and two holdouts.
C7	Baseline replay report with stable hash	A complete replay changes the same semantic fingerprint	Baseline snapshot only; not universal idempotence.
C8	Interface diagram and authority ordering	APFC bypasses the independent safety controller	Design proposal; no controlled safety experiment.
C9	Demonstration video and specified command–body–effect protocol	Self-observation fails to improve sealed future behaviour or survives self-model ablation	Physical demonstration only; controlled learning evidence open.
C10	Cortex shell, parity tests, App Server events, workspace binding	Native input is silently sent to the model, observations are durably promoted without a gate, or clients fork the owned thread	Depends on host terminal facilities, App Server behaviour, and tmux.
C11	Contextual fixture report with matched wind signal	Integrated context fails to improve the declared choice or fails to select different actions	Two authored cases; no phenomenal or open-domain inference.

Table 6. Observed values, separated by repository state.

Field	State	Value
Canonical repository	both	https://github.com/ciaoidea/MD-OS
Evaluated commit	B0	8e988b7f3fa677993afa5040ff8534002b1bc83e
Package / license	B0	5.0.1 / GPL-2.0-only
Node.js / host	B0	24.19.0 / Linux 6.18.35 x86-64
Core test result	B0	194 pass; 0 fail; 0 skip
Declared build	B0	pass
Verifier fixture	B0	1 verified; 2 rejected; 368 ms
Learning holdout	B0	0/2 before; 2/2 after; equal one-attempt budgets
Hypothesis reduction	B0	16 to 1; 4 bits
Replay hash	B0	d6d57a8196acde3f80fd30dd51b7f1743f93c0fca61ba48c1d93360076270f09
Post-replay learning counts	B0	2 episodes; 1 historical promotion; 0 runtime-eligible promotions
Integration commit	B1	02b099082e787ac1db26a0560857df7726beec67
Revision tests	B1	220/220 Node; 41/41 shell parity
Revision replay status	B1	APFC and commitment gate ok; compiler and aggregate health critical; two requested passes did not report matched_before: true
Revision replay hash	B1	2b5f7134b1ecd425fce75475baea3850c15941f98ef223499363911a48dd1d
Revision learning counts	B1	3 episodes; 0 promoted; 0 runtime-eligible promotions
Contextual fixture	B1	signal-only 1/2; integrated context 2/2; accuracy delta 0.5
Release health	both	runtime operable; active local workspace not publishable

- [19] OpenClaw Contributors. Agent runtime: Workspace contract and session bootstrap. GitHub documentation, 2026. URL <https://github.com/openclaw/openclaw/blob/main/docs/concepts/agent.md>. Accessed 15 August 2026.
- [20] OpenClaw Contributors. Openclaw documentation: Self-hosted gateway for agentic assistants. Online documentation, 2026. URL <https://docs.openclaw.ai/>. Accessed 15 August 2026.
- [21] OpenClaw Contributors. Tools, skills, and plugins overview. GitHub documentation, 2026. URL <https://github.com/openclaw/openclaw/blob/main/docs/tools/index.md>. Accessed 15 August 2026.
- [22] J. Kevin O’Regan and Alva No’e. A sensorimotor account of vision and visual consciousness. *Behavioral and Brain Sciences*, 24(5):939–973, 2001. doi: 10.1017/S0140525X01000115. URL <https://pubmed.ncbi.nlm.nih.gov/12239892/>.
- [23] Charles Packer, Vivian Fang, Shishir G. Patil, Kevin Lin, Sarah Wooders, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*, 2023. URL <https://arxiv.org/abs/2310.08560>.
- [24] Joon Sung Park, Joseph C. O’Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023. doi: 10.1145/3586183.3606763. URL <https://arxiv.org/abs/2304.03442>.
- [25] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/1b44b878bb782e6954cd888628510e90-Abstract-Conference.html.
- [26] Weinan Sun, Madhu Advani, Nelson Spruston, Andrew M. Saxe, and James E. Fitzgerald. Organizing memories for generalization in complementary learning systems. *Nature Neuroscience*, 26:

- 1438–1448, 2023. doi: 10.1038/s41593-023-01382-9. URL <https://www.nature.com/articles/s41593-023-01382-9>.
- [27] Martin Ulrich, Johannes Keller, Klaus Hoenig, Christian Waller, and Georg Gr"on. Neural correlates of experimentally induced flow experiences. *NeuroImage*, 86:194–202, 2014. doi: 10.1016/j.neuroimage.2013.08.019. URL <https://pubmed.ncbi.nlm.nih.gov/23959200/>.
- [28] Gido M. van de Ven, Hava T. Siegelmann, and Andreas S. Tolias. Brain-inspired replay for continual learning with artificial neural networks. *Nature Communications*, 11:4069, 2020. doi: 10.1038/s41467-020-17866-2. URL <https://www.nature.com/articles/s41467-020-17866-2>.
- [29] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023. URL <https://arxiv.org/abs/2305.16291>.
- [30] Daniel M. Wolpert, Zoubin Ghahramani, and Michael I. Jordan. An internal model for sensorimotor integration. *Science*, 269 (5232):1880–1882, 1995. doi: 10.1126/science.7569931. URL <https://doi.org/10.1126/science.7569931>.
- [31] Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Tuo Hua, Zhoujun Cheng, Dongchan Shin, Fangyu Lei, Yitao Liu, Yiheng Xu, Shuyan Zhou, Silvio Savarese, Caiming Xiong, and Victor Zhong. OSWorld: Benchmarking multimodal agents for open-ended tasks in real computer environments. *arXiv preprint arXiv:2404.07972*, 2024. URL <https://arxiv.org/abs/2404.07972>.
- [32] John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. SWE-agent: Agent-computer interfaces enable automated software engineering. *arXiv preprint arXiv:2405.15793*, 2024. URL <https://arxiv.org/abs/2405.15793>.
- [33] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2210.03629>.